

Additional Readings for XCS224U

Domain adaptation for supervised sentiment

- [Recursive Deep Models for Semantic Compositionality Over a Sentiment Treebank](#)
- [DynaSent: A Dynamic Benchmark for Sentiment Analysis](#)
- [Contextual Word Representations: A Contextual Introduction](#)
- [The Pile: An 800GB Dataset of Diverse Text for Language Modeling](#)
- [Advances in Neural Information Processing Systems 30 \(NIPS 2017\)](#)
- [The Annotated Transformer](#)
- [Self-Attention with Relative Position Representations](#)
- [Improving Language Understanding by Generative Pre-Training](#)
- [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#)
- [RoBERTa: A Robustly Optimized BERT Pretraining Approach](#)
- [ELECTRA: PRE-TRAINING TEXT ENCODERS AS DISCRIMINATORS RATHER THAN GENERATORS](#)
- [Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer](#)
- [BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension](#)
- [DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter](#)
- [Diffusion-LM Improves Controllable Text Generation](#)

Retrieval augmented in-context learning

- [Building Scalable, Explainable, and Adaptive NLP Models with Retrieval](#)
- [ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT](#)
- [Dense Passage Retrieval for Open-Domain Question Answering](#)
- [SPLADE: Sparse Lexical and Expansion Model for First Stage Ranking](#)
- [Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks](#)
- [Language Models are Few-Shot Learners](#)
- [Training language models to follow instructions with human feedback](#)
- [Chain-of-Thought Prompting Elicits Reasoning in Large Language Models](#)

- [Internet-augmented language models through few-shot prompting for open-domain question answering](#)
- [Demonstrate-Search-Predict: Composing retrieval and language models for knowledge-intensive NLP](#)

Advanced behavioral evaluation

- [Adversarial Examples for Evaluating Reading Comprehension Systems](#)
- [Breaking NLI Systems with Sentences that Require Simple Lexical Inferences](#)
- [Inoculation by Fine-Tuning: A Method for Analyzing Challenge Datasets](#)
- [Stress Test Evaluation for Natural Language Inference](#)
- [Adversarial NLI: A New Benchmark for Natural Language Understanding](#)
- [Dynabench: Rethinking Benchmarking in NLP](#)
- [COGS: A Compositional Generalization Challenge Based on Semantic Interpretation](#)
- [ReCOGS: How Incidental Details of a Logical Form Overshadow an Evaluation of Semantic Interpretation](#)

Analysis methods

- ["Why Should I Trust You?": Explaining the Predictions of Any Classifier](#)
- [BERT Rediscovered the Classical NLP Pipeline](#)
- [Axiomatic Attribution for Deep Networks](#)
- [Faithful, Interpretable Model Explanations via Causal Abstraction](#)
- [Inducing Causal Structure for Interpretable Neural Networks](#)
- [Finding Alignments Between Interpretable Causal Variables and Distributed Neural Representations](#)
- [Thread: Circuits What can we learn if we invest heavily in reverse engineering a single neural network?](#)

NLP methods

- [The Handbook of Computational Linguistics and Natural Language Processing](#)
- [Linguistic Structure Prediction by Noah A. Smith](#)
- [Dynaboard: An Evaluation-As-A-Service Platform for Holistic Next-Generation Benchmarking](#)
- [Moving Beyond Downstream Task Accuracy for Information Retrieval Benchmarking](#)

Your projects

- [Advice for Research Students](#)
- [Stuart Shieber on reporting research results](#)
- [Novelist Cormac McCarthy's tips on how to write a great science paper](#)
- [Geoff Pullum's Five Golden Rules](#)
- [Patrick Blackburn: How to give a good talk](#)
- [Datasheets for Datasets](#)
- [Model Cards for Model Reporting](#)